




# Networks (2IRR20)

Network Layer – Data Plane (04)

Dr. Tanir Ozcelebi

# This slide set


- Network layer key functions
- Data plane: Inside a router (forwarding)
- Internet Protocol (IP)



# Network layer key functions

# Router

- examines header fields in all IP datagrams passing through it
- moves each datagram in direction of destinations (to the corresponding output port)



# Two key network-layer functions

## network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

## analogy: taking a trip

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding




routing

# Network layer: data plane, control plane

## Data plane:

- *local* logic
- determines per-router forwarding behavior




## Control plane

- *network-wide* logic
- determines how datagram is routed end-to-end from source host to destination host
- 2 control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers


# Per-router control plane

Components of a distributed routing algorithm (*in each and every router*) interact/coordinate in the control plane.



# Software-Defined Networking (SDN) control plane


Remote controller computes, installs forwarding tables in routers.




# Data plane: Inside a router (forwarding)

# Router architecture overview

high-level view of generic router architecture:




# Input port functions



- using header field values, lookup output port using forwarding table in input port memory (“*match plus action*”)
- **input port queuing:** if datagrams arrive faster than forwarding rate into switch fabric

# Input port functions



- using header field values, lookup output port using forwarding table in input port memory ("*match plus action*")
- **destination-based forw.:** based on destination IP (traditional)
- **generalized forw.:** forward based on any set of header fields

# Destination-based forwarding

<i>forwarding table</i>	
Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

# Longest prefix matching

longest prefix match

choose *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001 which interface?

11001000 00010111 00011000 10101010 which interface?

# Longest prefix matching

longest prefix match

choose *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 000111** *****	2
otherwise	3

match!

examples:

11001000 00010111 00010110 10100001 which interface?  
11001000 00010111 00011000 10101010 which interface?

# Longest prefix matching

longest prefix match

choose *longest* address prefix that matches destination address.

Destination Address Range					Link interface
11001000	00010111	00010***	*****	*	0
11001000	00010111	00011000	*****	*	1
11001000	00010111	00011***	*****	*	2
otherwise					3

(shorter) match!



11001uuu uuuu1111 00010110 10100001

which interface?

11001000 00010111 00011000 10101010

which interface?

examples:

# Longest prefix matching

longest prefix match


choose *longest* address prefix that matches destination address.

Destination Address Range					Link interface
11001000	00010111	00010***	*****	0	
11001000	00010111	00011000	*****	1	
11001000	00010111	00011***	*****	2	
other	longest prefix match!			<input checked="" type="checkbox"/>	3
11001000	00010111	00010110	10100001	which interface?	
11001000	00010111	00011000	10101010	which interface?	

examples:


# Switching fabrics

- transfer packet from input link to appropriate output link




# Input port queuing

- If switch fabric slower than input ports combined -> queueing may occur at input queues
  - *queueing delay and loss due to input buffer overflow!*
- Head-of-the-Line (HOL) blocking: queued datagram at front of queue prevents others in queue from moving forward




output port contention: only one red datagram can be transferred -- lower red packet is *blocked*



one packet time later: green packet experiences HOL blocking

# Output port queuing




- *Buffering* required when datagrams arrive from fabric faster than link transmission rate.

Datagrams can be lost due to congestion, lack of buffers.

- *Scheduling* chooses among queued datagrams for transmission.


Priority scheduling – who gets best performance, network neutrality.

# Output port queuing



- buffering when arrival rate via switch exceeds output line speed
  - *queueing (delay) and loss due to output port buffer overflow!*
  - *drop policy:* which datagrams to drop if no free buffers?

# (Output) buffer management




- **drop policy:** decides which packet to drop when packet arrives to a full buffer.

- **tail drop:** drop newly arriving packet

- **priority:** drop/remove on priority basis

Abstraction: queue



# Packet scheduling


Router also needs a *scheduling policy* for deciding which packet to send next on link.

- first come, first served (FCFS)
- priority based
- round robin
- weighted fair queueing

# Scheduling policies: FCFS

*FCFS scheduling:*


- packets transmitted in order of arrival to output port
- a.k.a. first in, first out (FIFO)



# Scheduling policies: priority

## *Priority based scheduling:*


- arriving traffic classified, queued by class
  - any header fields can be used for classification
- send packet from highest priority queue that has buffered packets
  - FCFS within priority class



# Scheduling policies: round robin

## *Round Robin (RR) scheduling:*

- arriving traffic classified, queued by class
  - any header fields can be used for classification
- server cyclically, repeatedly scans class queues, sending one complete packet from each class (if available) in turn




# Scheduling policies: weighted fair queueing

## *Weighted Fair Queueing (WFQ):*

- generalized Round Robin
- each class,  $i$ , has weight,  $w_i$ , and gets weighted amount of service in each cycle:

$$\frac{w_i}{\sum_j w_j}$$


- minimum bandwidth guarantee (per-traffic-class)




# Internet Protocol (IP)

# Network Layer: Internet

host, router network layer functions:




# IP Datagram format



# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 = 


11011111	00000001	00000001	00000001
----------	----------	----------	----------

# IP addressing: introduction

Q: how are interfaces actually connected?

A: wired  
Ethernet interfaces connected by Ethernet switches

*For now:* don't need to worry about how one interface is connected to another (with no intervening router) – *a topic for our Link Layer lectures.*



A: wireless WiFi interfaces connected by WiFi base station


# Subnets

## ■ What's a subnet ?

- device interfaces that can physically reach each other **without passing through an intervening router**

## ■ IP addresses have structure:

- subnet part:** devices in same subnet have common high order bits
- host part:** remaining low order bits




network consisting of 3 subnets

# Subnets


*Recipe for defining subnets:*

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*



# Subnets


- where are the subnets?
- what are the /24 subnet addresses?



# IP addressing: CIDR

CIDR: Classless Inter-Domain Routing (pronounced “cider”)

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



# IP addresses: how to get one?

- can be hard-coded in config file (e.g., /etc/rc.config in UNIX)

OR

- **DHCP: Dynamic Host Configuration Protocol**
  - dynamically get address from server (“plug-and-play”)

# DHCP: Dynamic Host Configuration Protocol


**goal:** host *dynamically* obtains IP address from network server when it “joins” network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/on)
- support for mobile users who join/leave network

## DHCP overview:

- host broadcasts **DHCP discover** msg [optional]
- DHCP server responds with **DHCP offer** msg [optional]
- host requests IP address: **DHCP request** msg
- DHCP server sends address: **DHCP ack** msg

# DHCP client-server scenario




Typically, DHCP server will be co-located in router, serving all subnets to which router is attached

arriving **DHCP client** needs address in this network


# DHCP client-server scenario

DHCP server: 223.1.2.5



# DHCP client-server scenario

DHCP server: 223.1.2.5



Arriving client


The two steps above can be skipped “if a client remembers and wishes to reuse a previously allocated network address” [RFC 2131]

# DHCP: more than IP addresses

DHCP can return more than just allocated IP address:


- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

# DHCP: example



- Connecting laptop will use DHCP to get IP address, address of first-hop router, address of DNS server.
- DHCP REQUEST message encapsulated in UDP, encapsulated in IP, encapsulated in Ethernet
- Ethernet frame broadcast (dest: FFFFFFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demux'ed to IP demux'ed, UDP demux'ed to DHCP

# DHCP: example



- DHCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulated DHCP server reply forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

# IP addresses: how to get one?

**Q:** how does *network* get subnet part of IP address?

**A:** gets allocated portion of its provider ISP's address space

ISP's block      11001000 00010111 00010000 00000000    200.23.16.0/20

ISP can then allocate out its address space in 8 blocks:

Organization 0    11001000 00010111 00010000 00000000    200.23.16.0/23

Organization 1    11001000 00010111 00010010 00000000    200.23.18.0/23

Organization 2    11001000 00010111 00010100 00000000    200.23.20.0/23

...

.....


.....

.....

Organization 7    11001000 00010111 00011110 00000000    200.23.30.0/23


# Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:




# Hierarchical addressing: more specific routes

- Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



# IP fragmentation/reassembly


- network links have MTU (max. transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at *destination*
  - IP header bits used to identify, order related fragments



# IP fragmentation/reassembly

## example:

- 4000 byte datagram
- MTU = 1500 bytes




# IPv6: motivation

- **initial motivation:** 32-bit IPv4 address space would be completely allocated
- additional motivation:
  - speed processing/forwarding: 40-byte fixed length header
  - enable different network-layer treatment of “flows”

# IPv6 datagram format

**priority:** identify priority among datagrams in flow

**128-bit**  
IPv6 addresses




**flow label:** identify datagrams in same "flow." (concept of "flow" not well defined).

What's missing (compared with IPv4):

- no checksum (to speed processing at routers)
- no fragmentation/reassembly
- no options (available as upper-layer, next-header protocol at router)


# Transition from IPv4 to IPv6

- not all routers can be upgraded simultaneously
  - no “flag days”
  - how will network operate with mixed IPv4 and IPv6 routers?
- **tunneling:** IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers (“packet within a packet”)




# Tunneling and encapsulation

Ethernet connecting  
two IPv6 routers:




IPv4 network  
connecting two  
IPv6 routers




# Tunneling and encapsulation


Ethernet connecting  
two IPv6 routers:



IPv4 tunnel  
connecting two  
IPv6 routers



# Tunneling



Note source and destination addresses!